

Imbalanced multi-label classification using multi-task learning with extractive summarization

John Brandt

Yale University / New Haven, CT 06511

john.brandt@yale.edu

Abstract

Extractive summarization and imbalanced multi-label classification often require vast amounts of training data to avoid overfitting. In situations where training data is expensive to generate, leveraging information between tasks is an attractive approach to increasing the amount of available information. This paper employs multi-task training of an extractive summarizer and an RNN-based classifier to improve summarization and classification accuracy by 50% and 75%, respectively, relative to RNN baselines. We hypothesize that concatenating sentence encodings based on document and class context increases generalizability for highly variable corpuses.¹

1 Introduction

Multi-label sentence classification is a common applied task in natural language processing. Downstream applications include the identification of themes and topics in documents, quantifying differences between documents, emotion detection, and document labelling. In this paper, we seek to address the application of identifying and tagging relevant sentences across many documents of the same domain. We build upon a recurrent neural network classifier by leveraging context information from the source documents by jointly training a surrogate extractive summarizer. It is hypothesized that concatenating sentence encoding with context information encoded by the summarization hidden layers parallels the natural decision making process undertaken by human classification, which uses some combination of semantics and context.

¹Code is available on the author's GitHub.

2 Methods

2.1 Model Architecture

The model architecture is comprised of an extractive summarizer and a sentence classifier (Figure 1). In the extractive summarizer, pre-trained GloVe embeddings (Pennington et al., 2014) are concatenated with a learned embedding layer and are fed sequentially to a bi-directional gated recurrent unit (GRU) layer (Cho et al., 2014). The bidirectional layer contains both a forward GRU, reading words sequentially from w_1 to w_n , and a reverse GRU, reading words sequentially in reverse from w_n to w_1 . Attention is applied over the words to encode the sentences using the equations in (1-3) proposed by Yang et al. (2016).

$$u_{it} = \tanh(W_w h_{it} + b_s), \quad (1)$$

$$\alpha_i = \frac{\exp(u_{it}^T u_w)}{\sum_t \exp(u_{it}^T u_w)}, \quad (2)$$

$$s_i = \sum_t \alpha_{it} h_{it} \quad (3)$$

Then, a sentence-level GRU incorporates cross-sentence context in a binary classification setting to identify whether the sentence is relevant. Here, relevance is determined as whether the sentence belongs to any of the classes.

The sentence classifier passes pre-trained GloVe embeddings to a bi-directional GRU layer. After applying attention over the hidden states, a sigmoid activation is used to classify multiple labels for each sentence. The encoding output from the sentence GRU in the document summarizer is shared with the output of the word-level GRU in the sentence classifier. This is done by passing a matrix identifying, for each sample, the position of the relevant encoding in the summarizer output matrix. A lambda layer is used to slice the

summarizer output and share the summarizer sentence representation with the classifier network as shown in (4), where h_e is the extractor GRU output, h_c is the classifier GRU output, and i is the index of the h_e matrix corresponding to h_c

$$(ih_{e[:,i,:]}) \oplus h_c \quad (4)$$

2.2 Training

Sentences were tokenized with a max word count of 10,000 and encoded with pre-trained GloVe embeddings. Contractions and punctuation were treated as their own word. Sentences were padded to 30 words. Training was done on a Linux cluster with TensorFlow, using the Adam optimizer with a learning rate of $2.5e - 5$. The sentence extractor GRU was regularized with an l1-l2 regularizer ($1e - 5$). Classes were weighted by their inverse frequency for both the extractor and the summarizer inputs. In order to maintain equal sample sizes for each input, documents were broadcasted to share dimensions with classified sentences. For each document and epoch, the extractor weights were thus updated n times, where n is the number of relevant sentences in the document. The loss function is the sum of the binary cross entropy loss for each output (5).

$$-\sum y_s(x) \log y_s(x) - \sum y_d \log y_d(x) \quad (5)$$

The architecture was initially built with a 60/20/20 training, validation, test split. All results presented in this paper use an 80/20 training test split.

2.3 Data

The data set was aggregated by Resource Watch (World Resources Institute, 2018). Sentence relevance to 17 separate classes were hand-coded for 151 national environmental policies, totalling approximately 4,000 pages, by a team of domain experts. More specifically, sentences and phrases were classified according to their relevance to the Sustainable Development Goals, a set of 17 goals adopted by the United Nations to guide global development agenda. Class labels are shown in Table 1. There is significant overlap between the classes, and many sentences and phrases are identified as belonging to multiple classes.

The original data set only included the extracted sentences and phrases, as well as the raw HTML

Class	Samples
No Poverty	15
Zero Hunger	850
Good Health	1042
Education	322
Gender Equality	367
Clean Water	841
Clean Energy	2258
Economic Growth	241
Infrastructure	579
Reduced Inequality	65
Sustainable Cities	591
Responsible Production	625
Climate Action	1106
Aquatic Life	283
Land Life	1056
Peace and Justice	121
Partnerships	406

Table 1: Number of samples per class. Note that the sum of the samples exceeds the sample size due to the multi-label nature of the data.

Model Type	Loss	Top-1	Top-3
GRU	0.2626	0.368	0.601
Multi-task	0.2065	0.553	0.791

Table 2: Classification metrics for baseline (GRU) and multi-task network.

Model Type	Loss	Precision	Recall
Hierarchical GRU	0.400	0.236	0.312
Multi-task	0.181	0.460	0.459

Table 3: Summarization metrics for baseline (GRU) and multi-task network.

files of the policies. Extracted phrases were matched to their location in source documents using a combination of regular expression and the R package stringR. Four of the the 155 original documents were not considered because of differences between source and extracted sentences. All of the 151 documents used were matched at 100% accuracy.

3 Experiments

To test the effectiveness of sharing weights between a sentence classifier and an extractive summarization network, an ablation study was used to test the multi-task network against its individual

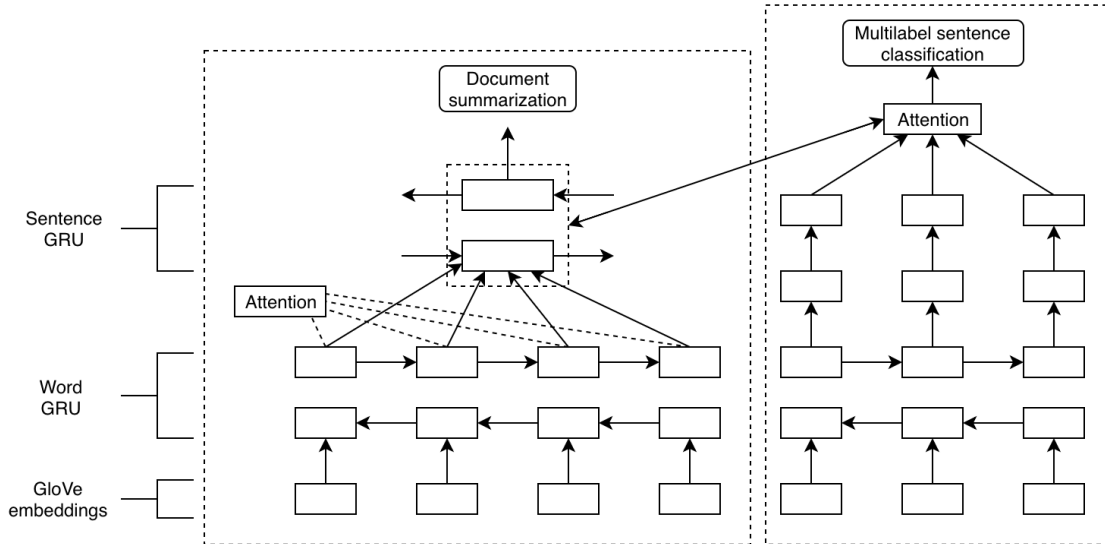


Figure 1: Model architecture comprised of document summarization (left) and multi-label sentence classification (right) indicating the shared layer and the position of attention layers.




Figure 2: Baseline GRU (green) and multi-task (purple) sentence classification accuracy by class demonstrates that, for each of the classes, leveraging the source document context improves accuracy.

components. Sentence classification was tested using a GRU with attention, and extractive summarization was tested with a hierarchical GRU

with attention over word embeddings.

Classification performance was evaluated by considering loss, accuracy at top-K, and accuracy for each class, while summarization performance was evaluated with precision and recall. All models were trained with the same training, validation, and test data, and the epoch with the lowest validation loss was used for comparison. Data was split to ensure that no documents or sentences were shared between validation and test data.

The multi-task model outperforms the RNN baseline in terms of top-1 accuracy, percent accuracy by class, and predicting the number of labels for a candidate sentence. As shown in Table 2, the multi-task model increases top-1 accuracy by 20% absolute over the RNN baseline. Figure 2 shows the accuracy by class, demonstrating that the multi-task model significantly improves upon the baseline for identifying extremely unbalanced classes.

For predicting the number of labels, the multi-task model improves upon the RNN baseline significantly. While the RNN baseline failed to identify any labels 23% of the time, the multi-task model reduces this error to 13%. As the confusion matrix in Figure 3 shows, the multi-task model also identifies the correct number of labels for 1 and 2 label sentences on average 40% more frequently than the baseline. Both the RNN baseline and the multi-task model perform poorly when identifying 3, 4, and 5 label sentences, though

these represent less than 1% of the dataset cumulatively.

Updating weights for extractive summarization based upon their contribution to sentence classification accuracy greatly improved the summarization performance over the hierarchical GRU baseline. As Table 3 shows, the multi-task model achieved 100% and 48% relative gains in precision and recall over the baseline.




Figure 3: Confusion matrix of number of topics identified by baseline and multi-task classifier indicates that the multi-task classifier predicts label number more accurately.

4 Discussion

The results suggest that, in applications where sentences are extracted from longer documents, jointly training an extractive summarization network effectively allows the classifier to leverage context information about the source documents to improve model accuracy. This is especially important for highly context-specific applications such as the present one, where the phrasing of policy

agenda is dependent upon the national infrastructure, political context, and development goals, as well as the translation quality between the source language and English.

While the extractive summarization is by no means state of the art, it allows the sentence classifier to learn and predict what makes a sentence relevant enough to be classified. This is a much different view of the sentences than that obtained by the sentence classifier, which simply seeks differences between classes rather than the underlying information that makes the sentence relevant.

References

- Kyunghyun Cho, Bart van Merriënboer, Çalar Gülçehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. 2014. [Learning phrase representations using rnn encoder-decoder for statistical machine translation](#). In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. Association for Computational Linguistics, Doha, Qatar, pages 1724–1734. <http://www.aclweb.org/anthology/D14-1179>.
- Jeffrey Pennington, Richard Socher, and Christopher D Manning. 2014. Glove: Global vectors for word representation. In *EMNLP*. volume 14, pages 1532–1543.
- World Resources Institute. 2018. [Resource watch](#). <http://resourcewatch.org/>.
- Zichao Yang, Diyi Yang, Chris Dyer, Xiaodong He, Alexander J. Smola, and Eduard H. Hovy. 2016. Hierarchical attention networks for document classification. In *HLT-NAACL*.